



Universität
Zürich UZH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

DaST 
Data • (Systems+Theory)



Computer Vision
and Learning
Group

Advancing (Sparse) Multi-View Depth Estimation

Bachelor's Thesis

Submission of the thesis

April 29, 2026

Fabio Di Meo
University of Zurich

Student-ID: 22-929-798
fabio.dimeo@uzh.ch

Supervisors:
Prof. Dr. Dan Olteanu,
Frano Rajič

Abstract

Zusammenfassung

Acknowledgements

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Contents

1	Introduction	1
1.1	Background and Motivation	1
1.2	Problem Statement and Contributions	2
2	Literature Review and Background	4
3	Methodology: The Evaluation Framework	5
4	System Implementation	6
5	Results	7
6	Refinement Strategy	8
7	Discussion	9
8	Conclusions	10
	Bibliography	10
	List of Acronyms	12
A	Artifacts Availability	13

1

Introduction

Recovering the three-dimensional structure of a scene from visual observations is a foundational problem in computer vision with critical applications in robotics and autonomous navigation [2]. While spatial-first foundation models such as MAST3R [4] and PI3X [7] have achieved remarkable per-frame geometric fidelity, their application to 4D sequences - where temporal coherence is as vital as spatial accuracy - remains an open challenge. In sparse-view configurations (2-4 cameras), the problem is bifurcated: spatial models lack the temporal grounding to prevent high-frequency jitter, whereas dedicated 4D models, such as VGGT4D [3], introduce temporal priors that risk over-smoothing geometry or creating "motion-blur" artifacts.

This thesis investigates the transition from per-frame 3D reconstruction to sequence-level 4D modeling. We evaluate whether the current generation of foundation models can maintain metric integrity across time or if the introduction of temporal constraints inevitably degrades geometric precision.

1.1 Background and Motivation

Classical multi-view reconstruction methods such as SfM [5] provide high rigour but struggle with the textureless regions and dynamic occlusions inherent in human-centric environments. Learning-based spatial models like DUST3R [6] and PI3 [7] have largely superseded these by regressing dense pointmaps directly, yet they treat video as a collection of independent snapshots. This independence results in significant temporal instability when

these models are deployed on robotic platforms requiring smooth trajectory estimation.

To address this, temporal-first architectures like VGGT4D [3] and MonoFusion [8] incorporate motion cues to encourage sequence-level consistency. However, the nature of this consistency - whether it acts as a beneficial geometric anchor or a destructive low-pass filter - is poorly understood in sparse-view settings. There is a critical need for a benchmarking framework that quantifies the "temporal cost" of these various architectural choices.

1.2 Problem Statement and Contributions

The central objective of this work is to characterize the performance of 3D and 4D foundation models under the strict hardware constraints of sparse-view (2–4 camera) setups. We focus on measuring the divergence between per-frame accuracy and sequence-level stability, identifying where hierarchical vs. graph-based alignment strategies succeed in mitigating drift.

The main contributions of this work are:

- **4D Benchmarking Framework.** A unified evaluation pipeline for sparse multi-view video that subjects six base models (MAST3R-SfM, VGGT, PI3, PI3X, VGGT4D, and MonoFusion) to identical data-loading and semantic masking conditions across the DEX-YCB, Hi4D, and 4D-DRESS datasets. Crucially, this framework evaluates the impact of applying the GGPT [1] refinement strategy across these architectures (excluding MonoFusion), providing a direct comparison between raw spatial outputs and geometry-grounded 4D sequences.
- **The Delta-Consistency Metric.** We formalize a novel quantitative measure to isolate the "temporal cost" of sequence alignment. This metric, $\Delta_{\text{consistency}}$ (see Equation (1.1)), captures the divergence in geometric fidelity when per-frame reconstructions are transformed from independent, local coordinate systems into a globally consistent 3D reference frame. This value exposes the spatial degradation caused by alignment drift and temporal constraints.

$$\Delta_{\text{consistency}} = \text{Chamfer}_{4\text{D}} - \text{Chamfer}_{3\text{D}} \quad (1.1)$$

- **Comparative Alignment Study.** A rigorous analysis of three temporal alignment strategies, highlighting a performance split: identifying

hierarchical binary-tree methods as superior for static scene completeness and graph-based optimization (PGO) as superior for dynamic accuracy and minimizing $\Delta_{\text{consistency}}$.

This section must include explicit references to the detailed method descriptions in later chapters (hierarchical alignment trees + PGO formulation).

- **Staticity-Informed Refinement.** An investigation into refinement strategies - such as staticity constraints applied to MAST3R-SfM - aimed at reducing high-frequency jitter without the over-smoothing seen in purely temporal architectures.

2

Literature Review and Background

3

Methodology: The Evaluation Framework

4

System Implementation

5

Results

6

Refinement Strategy

7

Discussion

8

Conclusions

Bibliography


- [1] Yutong Chen *et al.* “GGPT: Geometry-Grounded Point Transformer”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2026.
- [2] R. I. Hartley and A. Zisserman. *Multiple View Geometry in Computer Vision*. Second. Cambridge University Press, ISBN: 0521540518, 2004.
- [3] Yu Hu *et al.* *VGGT4D: Mining Motion Cues in Visual Geometry Transformers for 4D Scene Reconstruction*. 2025. arXiv: [2511.19971](https://arxiv.org/abs/2511.19971) [cs.CV]. URL: <https://arxiv.org/abs/2511.19971>.
- [4] Vincent Leroy, Yohann Cabon, and Jerome Revaud. *Grounding Image Matching in 3D with MAST3R*. 2024. arXiv: [2406.09756](https://arxiv.org/abs/2406.09756) [cs.CV].
- [5] Johannes Lutz Schönberger and Jan-Michael Frahm. “Structure-from-Motion Revisited”. In: *Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016.
- [6] Shuzhe Wang *et al.* “DUST3R: Geometric 3D Vision Made Easy”. In: *CVPR*. 2024.
- [7] Zihan Wang *et al.* “MonoFusion: Sparse-View 4D Reconstruction via Monocular Fusion”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. Oct. 2025, pp. 8252–8263.
- [8] Zihan Wang *et al.* “MonoFusion: Sparse-View 4D Reconstruction via Monocular Fusion”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. Oct. 2025, pp. 8252–8263.

List of Acronyms



Artifacts Availability

The artifacts associated with this thesis can be found in the ifi-GitLab¹. Within this repository, there are two main folders: one containing the code to fine-tune the model, and the other containing the code to generate the training data, including all the different datasets. The README file provided in the repository contains step-by-step instructions on how to set up the environment used in this thesis, install all necessary dependencies, and run the source code.



Upload your Code to Git-Lab rather than attaching it to the PDF.

¹<https://gitlab.ifi.uzh.ch/ddis/Students/Theses/<NAME>>



Universität
Zürich^{UZH}

Declaration of independence for written work

I hereby declare that I have composed this work independently and without the use of any aids other than those declared (including generative AI such as ChatGPT). I am aware that I take full responsibility for the scientific character of the submitted text myself, even if AI aids were used and declared (after written confirmation by the supervising professor). All passages taken verbatim or in sense from published or unpublished writings are identified as such. The work has not yet been submitted in the same or similar form or in excerpts as part of another examination.

Important: Update this text with the latest **Declaration of independence for written work** from <https://www.ifi.uzh.ch/en/studies/msc-info.html>. Edit me in `resources/declaration_en.tex`

Zürich, April 29, 2026

Signature of
Fabio Di Meo