



Universität
Zürich UZH

Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

DaST 
Data • (Systems+Theory)



VLG | Computer Vision
and Learning
Group

Advancing (Sparse) Multi-View Depth Estimation

Bachelor's Thesis

Submission of the thesis

June 9, 2026

Fabio Di Meo
University of Zurich

Student-ID: 22-929-798
fabio.dimeo@uzh.ch

Supervisors:
Prof. Dr. Dan Olteanu,
Frano Rajič

Abstract

Zusammenfassung

Acknowledgements

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Contents

1	Introduction	1
1.1	Background and Motivation	2
1.2	Problem Statement and Contributions	3
2	Literature Review and Background	6
2.1	Structure-from-Motion and Multi-View Stereo Foundations . .	6
2.2	Foundation Models for 3D and 4D Reconstruction	8
2.2.1	Spatial Reconstruction Models	8
2.2.2	Geometry Refinement Models	9
2.2.3	Temporal Reconstruction Models	10
2.3	SAM2 for Static and Dynamic Scene Segmentation	11
2.4	Mathematical Background	12
2.4.1	Similarity Transformations and Umeyama Alignment .	12
2.4.2	Rotation Geometry on $SO(3)$	13
2.4.3	Geodesic Rotation Averaging	13
2.4.4	Chamfer Distance	14
3	Methodology: The Evaluation Framework	15
4	System Implementation	16
5	Results	17
6	Refinement Strategy	18
7	Discussion	19
8	Conclusions	20
	Bibliography	20
	List of Acronyms	24

A Artifacts Availability**25**

1

Introduction

Recovering the three-dimensional structure of a dynamic scene from visual observations is a foundational challenge in computer vision with direct implications for robotics and autonomous navigation [9]. Recent spatial foundation models—such as MAST3R-SfM [13], VGGT [21], PI3, and PI3X [23]—have achieved remarkable per-frame geometric fidelity by regressing dense 3D pointmaps directly from images. Yet in 4D settings, where temporal coherence across a video sequence is as critical as spatial accuracy, these models fail: assembled frame-by-frame, their outputs exhibit severe high-frequency jitter that renders them unsuitable for robotics applications requiring smooth, metric-consistent depth. Geometry-refinement approaches such as GGPT [4] attempt to correct these inconsistencies by grounding spatial predictions with sparse SfM guidance, but fail to generalise to sparse-view dynamic scenes, worsening Chamfer distance by 3–10 $\times$. Conversely, dedicated 4D architectures—VGGT4D [11] and MonoFusion [24]—sacrifice geometric precision to buy temporal smoothness, inflating Chamfer distance by 20–160 % and producing systematic alignment failures in sparse-view settings.

The central finding of this thesis is that a *decoupled* paradigm outperforms both families: take a strong per-frame 3D estimator, align its outputs into a common coordinate frame through a lightweight temporal alignment strategy, and apply a confidence-weighted, mask-aware Gaussian post-processing filter (Opt) to enforce static-scene consistency. This approach reduces jitter by 25–81 % across four models (MASt3R-SfM, VGGT, PI3, PI3X) while preserving geometry to within ± 2 –8 % Chamfer distance—a strictly superior trade-off to

every native 4D architecture and geometry-refinement baseline we evaluate. At the two-camera extreme, VGGT combined with Opt beats the native 4D architecture VGGT4D on every single metric.

1.1 Background and Motivation

Classical multi-view reconstruction pipelines such as Structure-from-Motion (SfM) [17] produce accurate, metric-scale geometry in well-constrained settings: they rely on repeatable keypoint matching and bundle adjustment, which deliver precise camera poses and sparse point clouds when images share sufficient texture, overlap, and baseline. In sparse-view configurations of two to four cameras, however, these conditions frequently fail. Low inter-camera overlap means that few keypoints are visible in more than one view simultaneously [9]; short or degenerate baselines reduce parallax and destabilise triangulation [17]; and textureless surfaces—common on human skin, clothing, and manipulation objects—yield too few reliable correspondences for bundle adjustment to converge [22, 13]. The result is either outright reconstruction failure or incomplete, drift-prone geometry that cannot support downstream tasks such as grasping or object tracking.

Learning-based spatial models such as DUS3R [22], PI3, and PI3X [23] overcome many of these failure modes by directly regressing dense 3D pointmaps from images, using learned priors to fill in low-texture and low-overlap regions where classical methods collapse. These models achieve strong per-frame accuracy—PI3 and PI3X in particular deliver the highest accuracy and completeness at a strict 1 cm robotics threshold among all models we evaluate. As purely spatial estimators, however, they process each frame independently and provide no explicit mechanism for enforcing temporal consistency across a sequence. The resulting frame-to-frame instability—manifesting as high-frequency jitter, scale fluctuations, or pose inconsistencies—can render otherwise accurate reconstructions unsuitable for downstream tasks including trajectory estimation, object tracking, and robotic manipulation.

Geometry-refinement approaches such as GGPT [4] take a different direction: rather than addressing temporal stability, they seek to improve the multi-view geometric consistency of per-frame feed-forward reconstructions by grounding dense pointmap predictions with sparse SfM-derived geometric guidance. Under the sparse-view dynamic conditions considered in this thesis, however, GGPT worsens Chamfer distance by 3–10 $\times$ across all evaluated backbone models, making it unsuitable for geometry-critical applications in this setting. Dedicated 4D architectures—VGGT4D [11] and Mono-Fusion [24]—enforce sequence-level consistency, but at a measurable cost:

VGGT4D inflates Chamfer distance by 20–160% relative to its per-frame baseline, and MonoFusion suffers systematic alignment convergence failures in sparse-view settings, producing duplicated subject artefacts rather than coherent reconstructions.

None of these approaches simultaneously achieves geometric fidelity and temporal consistency under sparse-view dynamic conditions. These results motivate the central design principle of this work: temporal consistency and geometric fidelity are not in fundamental conflict—they simply require a targeted rather than global smoothing strategy.

1.2 Problem Statement and Contributions

We address the problem of temporally consistent, geometrically precise 4D reconstruction from sparse multi-view video (two to four cameras) of dynamic scenes, evaluated on the DEX-YCB [3] and Hi4D [26] datasets. The main contributions of this work are:

- **Unified 4D Benchmarking Framework.** A structured evaluation pipeline that subjects six base models—MASt3R-SfM, VGGT, PI3, PI3X, VGGT4D, and MonoFusion—to identical data-loading, semantic masking, and metric-computation conditions across DEX-YCB [3] and Hi4D [26]. The framework measures both 3D per-frame and 4D sequence-level performance via Chamfer distance, accuracy and completeness at a 1 cm threshold, temporal jitter, drift, and camera pose error, with static/dynamic decomposition where applicable (see Chapter 3). It additionally evaluates the effect of GGPT [4] geometry refinement applied on top of all base models except MonoFusion.
- **Comparative Temporal Alignment Study.** An analysis of three strategies for lifting per-frame 3D reconstructions into a globally consistent 4D representation (see Chapter 3). The Reference strategy aligns every frame to a common anchor using static-scene correspondences and the Umeyama similarity transform. The Hierarchical strategy builds a binary tree of pairwise alignments, propagating consistency bottom-up. The Pose-Graph Optimisation strategy computes all $\mathcal{O}(T^2)$ pairwise edges, weights them by alignment reliability, and refines all frame poses simultaneously via robust geodesic averaging on $\text{SO}(3)$. The Hierarchical strategy consistently delivers the best static-scene completeness and geometric stability; Pose-Graph Optimisation minimises the temporal consistency cost $\Delta_{\text{consistency}}$ (Equation 1.1) and best preserves dynamic-region accuracy.

$$\Delta_{\text{consistency}} = \text{Chamfer}_{4\text{D}} - \text{Chamfer}_{3\text{D}} \quad (1.1)$$

This quantity isolates the spatial degradation introduced purely by coordinate alignment, enabling direct comparison of alignment strategies and model families.

- Confidence-Weighted Temporal Smoothing (Opt).** A ground-truth-free post-processing method that reduces frame-to-frame instability for any per-frame estimator without retraining (see ??). Opt applies a Gaussian temporal filter to already-aligned pointmaps, weighted by model confidence and restricted to pixels classified as static by SAM2 [16]-derived masks—leaving dynamic pixels untouched to avoid suppressing real motion. A Pareto-frontier sweep over temporal bandwidth $\sigma \in \{0.0, 0.5, 1, 2, 4, 6, 8, 12, 20, 30\}$ and blending strength $\alpha \in \{0.0, 0.2, 0.5, 0.8, 1.0\}$ demonstrates that the optimal configuration is model-dependent: MAST3R-SfM benefits from a wider temporal window ($\sigma=6$, $\alpha=1.0$), PI3 and PI3X from a moderate window ($\sigma=4$, $\alpha=1.0$), and VGGT from a narrower window ($\sigma=2$, $\alpha=1.0$), reflecting distinct jitter profiles across architectures. These per-model configurations achieve a 25–81 % jitter reduction while preserving geometry to within ± 2 –8 % Chamfer distance—a strictly superior trade-off to every native 4D architecture and geometry-refinement baseline evaluated in this thesis.
- Empirical Characterisation of Architectural Families.** A systematic comparison of spatial-first models (PI3, PI3X, MAST3R-SfM, VGGT), the geometry-refinement approach GGPT, and native 4D architectures (VGGT4D, MonoFusion) reveals distinct and quantifiable failure modes under sparse-view dynamic conditions. PI3 and PI3X consistently achieve the strongest geometric performance across both DEX-YCB and Hi4D. MAST3R-SfM exhibits the greatest frame-to-frame instability, making it the model that benefits most from temporal refinement. VGGT provides a competitive balance between accuracy and stability. GGPT worsens Chamfer distance by 3–10 $\times$ and produces inconsistent accuracy and completeness behaviour across architectures and view counts, while suppressing jitter to near-zero—a combination that makes it unsuitable for geometry-critical applications despite its apparent temporal stability. Native 4D architectures reduce temporal noise but consistently sacrifice geometric accuracy relative to spatial-first models. Together these findings establish the decoupled framework—strong spatial estimation, robust temporal alignment,

and confidence-weighted temporal smoothing—as the most effective approach for sparse-view 4D reconstruction.

2

Literature Review and Background

This chapter reviews the theoretical and technical foundations underlying the reconstruction pipelines evaluated throughout this thesis. Section 2.1 introduces classical Structure-from-Motion (SfM) and Multi-View Stereo (MVS) methods, highlighting both their strengths and their limitations under sparse-view conditions. Section 2.2 surveys recent foundation models for 3D vision, grouped according to their treatment of spatial and temporal information. Section 2.3 describes the Segment Anything Model 2 (SAM2), which is used throughout this work to distinguish static and dynamic scene regions. Finally, Section 2.4 presents the mathematical concepts required for the proposed alignment and temporal smoothing framework, including similarity transformations, rotation averaging on $SO(3)$, robust geodesic averaging, and geometric evaluation metrics.

2.1 Structure-from-Motion and Multi-View Stereo Foundations

Three-dimensional scene reconstruction from multiple images has traditionally been addressed through the complementary paradigms of Structure-from-Motion (SfM) and Multi-View Stereo (MVS). Together, these methods form the foundation of modern photogrammetry systems and remain the dominant approach for producing accurate metric-scale reconstructions from calibrated or partially calibrated image collections [9, 17, 18].

SfM seeks to jointly estimate camera poses and scene structure from im-

age observations. Given a collection of overlapping images, distinctive visual features are first detected and matched across views. Geometric verification is then used to reject inconsistent correspondences, after which camera poses and sparse three-dimensional points are estimated through triangulation and refined using bundle adjustment [9, 17]. Modern SfM systems such as COLMAP [17] achieve remarkable accuracy and robustness on large-scale reconstruction tasks, often producing camera trajectories with sub-pixel reprojection error.

While SfM recovers only a sparse representation of scene geometry, MVS extends this sparse reconstruction into a dense surface estimate. Given camera poses obtained from SfM, MVS algorithms estimate depth for individual pixels by enforcing photometric consistency across multiple views. These depth estimates are subsequently fused into dense point clouds or surface representations [18, 19]. Modern MVS systems employ sophisticated view-selection strategies and regularisation techniques that enable high-quality reconstructions across a wide range of scenes [18].

Despite their success, classical SfM and MVS pipelines rely fundamentally on the existence of reliable multi-view correspondences. Performance degrades substantially when these assumptions are violated. Sparse-view reconstruction represents a particularly challenging regime because the number of observations available for establishing geometric constraints is severely limited. When only two to four cameras are available, several failure modes become prominent.

First, low inter-camera overlap reduces the number of features that are simultaneously visible across views, limiting the quantity and quality of geometric correspondences available for pose estimation [9, 15]. Second, short or degenerate baselines reduce triangulation accuracy because depth uncertainty grows rapidly as the triangulation angle decreases [9, 17]. Third, textureless surfaces such as human skin, clothing, walls, and many manipulation objects provide few distinctive features for matching, causing correspondence estimation and bundle adjustment to become unstable [22, 13]. Dynamic scene content introduces an additional challenge because moving objects violate the static-world assumptions underlying both feature matching and bundle adjustment [7].

These limitations are particularly acute in robotics and human-centred perception scenarios, where sparse multi-camera systems are often preferred for reasons of cost, deployment complexity, and occlusion constraints, yet such configurations provide exactly the conditions under which traditional SfM and MVS become least reliable. The resulting reconstructions may be incomplete, geometrically inconsistent, or entirely unsuccessful, limiting their usefulness for downstream tasks such as object tracking, trajectory

estimation, manipulation planning, and scene understanding.

2.2 Foundation Models for 3D and 4D Reconstruction

Rather than relying on geometric matching alone, modern foundation models learn strong scene priors from large-scale image collections, enabling them to recover plausible geometry even in regions with limited texture, weak overlap, or ambiguous correspondence structure [22, 13, 21]. Although these methods share a common goal of reconstructing scene structure from images, they differ substantially in how geometry and time are represented.

The models evaluated in this thesis are grouped into three architectural families according to their primary design objective. *Spatial reconstruction models* estimate scene geometry from a collection of images without explicit temporal modelling (MASt3R-SfM, VGGT, PI3, and PI3X). *Geometry refinement models* improve existing reconstructions through additional geometric reasoning (GGPT). *Temporal reconstruction models* jointly model geometry and motion to produce temporally coherent reconstructions (VGGT4D and MonoFusion). Together, these families represent the principal design directions currently being explored in modern 3D and 4D vision, and their differing assumptions about geometry, correspondence, and temporal consistency form the basis for the comparative evaluation presented in Chapters 3 and 5.

2.2.1 Spatial Reconstruction Models

Spatial reconstruction models focus on recovering accurate scene geometry from a set of observations without explicitly modelling temporal evolution. Their primary objective is geometric reconstruction, with camera poses, depth, or pointmaps inferred directly from one or more images using learned geometric priors. While these models can be applied independently to each frame of a video sequence, they provide no mechanism for enforcing temporal consistency across successive reconstructions.

MASt3R-SfM. MASt3R-SfM is a hybrid reconstruction pipeline that combines learned geometric reasoning with classical SfM optimisation [8]. The approach builds upon DUST3R, which introduced dense pointmap regression as an alternative to traditional depth estimation and triangulation [22], and MASt3R, which further improved robustness by grounding image match-

ing directly in predicted three-dimensional geometry [13]. MAST3R-SfM integrates these learned representations into a complete SfM pipeline, using geometry-aware correspondences to initialise camera poses and scene structure before applying global optimisation. By combining learned matching with bundle adjustment, the system achieves state-of-the-art performance on challenging unconstrained reconstruction benchmarks while retaining the geometric consistency advantages of classical SfM.

VGGT. VGGT (Visual Geometry Grounded Transformer) represents a further step toward fully end-to-end geometric reasoning [21]. Rather than predicting isolated pairwise reconstructions, VGGT processes multiple views jointly using a transformer architecture trained to recover cameras, depth, and scene structure within a unified framework. By learning global geometric relationships directly from large-scale data, VGGT reduces dependence on handcrafted reconstruction stages and achieves strong performance across a wide variety of reconstruction tasks. Its architecture reflects a broader trend toward general-purpose geometric foundation models capable of solving multiple three-dimensional vision problems within a single learned representation.

PI3 and PI3X. PI3 introduces a reference-free formulation of visual geometry reconstruction that removes the need for a designated anchor image [23]. Using a fully permutation-equivariant architecture, the model jointly predicts camera poses and dense point maps from an unordered set of images, making reconstruction inherently robust to input ordering and reference-view selection. By learning geometric priors from large-scale data, PI3 can recover accurate scene structure even in sparse-view settings where traditional correspondence-based methods become unreliable. PI3X extends this framework with improved point-map decoding, more reliable confidence estimation, optional conditioning on camera and depth priors, and support for approximate metric-scale reconstruction. These modifications improve reconstruction quality while preserving the reference-free geometric reasoning of the original model.

2.2.2 Geometry Refinement Models

A complementary class of methods seeks to improve reconstruction quality without replacing the underlying reconstruction backbone. Instead of generating geometry directly, these approaches operate as post-processing stages that refine an existing geometric prediction using additional evidence.

GGPT. GGPT (Geometry-Grounded Point Transformer) is a geometry-refinement framework designed to enhance feed-forward three-dimensional reconstructions through geometry-guided correction [4]. The method combines dense pointmap predictions with sparse geometric cues derived from an auxiliary SfM process, using a transformer-based refinement network to improve consistency and reconstruction quality. Conceptually, GGPT occupies a middle ground between classical optimisation-based reconstruction and purely feed-forward prediction: geometric structure is first estimated by a backbone model and subsequently refined using additional geometric evidence.

Unlike spatial reconstruction and temporal reconstruction architectures, GGPT does not define a complete reconstruction pipeline in its own right. Instead, it functions as a model-agnostic refinement stage that can be applied to the outputs of existing reconstruction systems. This design makes it particularly relevant for understanding whether reconstruction quality is best improved through stronger backbone architectures, explicit temporal modelling, or post-hoc geometric refinement.

2.2.3 Temporal Reconstruction Models

A common limitation of purely spatial reconstruction models is the absence of explicit temporal reasoning [11]: each frame or image set is reconstructed independently, meaning that geometric predictions may vary between consecutive frames even when the underlying scene remains static. Temporal reconstruction models address this by incorporating motion information and sequence-level reasoning during inference, jointly modelling geometry and temporal evolution rather than treating each frame as an isolated input [11].

VGGT4D. VGGT4D extends the VGGT architecture to four-dimensional reconstruction by incorporating temporal motion cues into the transformer backbone [11]. The model is designed to reason jointly about spatial structure and temporal evolution, enabling reconstruction of dynamic scenes across time. By leveraging information from neighbouring frames, VGGT4D aims to improve temporal consistency and reduce frame-to-frame fluctuations that commonly arise in per-frame reconstruction systems.

MonoFusion. MonoFusion approaches sparse-view four-dimensional reconstruction from a different perspective [24]. Instead of relying solely on multi-view geometric constraints, it combines monocular depth priors with temporal fusion mechanisms to accumulate information across a sequence. This

strategy enables reconstruction from extremely sparse camera configurations and provides an explicit mechanism for enforcing temporal coherence. The model reflects a growing trend toward integrating learned monocular geometry with multi-view reasoning to overcome the limitations of sparse observations [14, 12, 5].

While temporal architectures address an important limitation of spatial models by introducing sequence-level consistency, the incorporation of temporal information inevitably introduces additional modelling assumptions and complexity. This creates a trade-off between geometric precision and temporal stability that remains an active area of research, and one that the comparative evaluation in this thesis aims to illuminate under controlled sparse-view conditions.

2.3 SAM2 for Static and Dynamic Scene Segmentation

SAM2 (Segment Anything Model 2) extends the original Segment Anything Model from static-image segmentation to image and video sequences [16]. By combining promptable segmentation with a memory mechanism that propagates object information across time, SAM2 produces temporally consistent masks throughout a video sequence. Trained on large-scale image and video datasets, the model demonstrates strong zero-shot generalisation across diverse object categories and scene types [16].

Video segmentation has become increasingly relevant for dynamic-scene reconstruction, where moving actors and objects violate the static-world assumptions commonly exploited by geometric reconstruction methods. Recent systems such as MonoFusion [24] employ foreground-background segmentation to distinguish dynamic scene content from temporally stable geometry during reconstruction and optimisation.

In this thesis, SAM2 is used to generate static and dynamic scene masks for evaluation and alignment purposes. These masks enable separate analysis of reconstruction quality in static and dynamic regions and provide the basis for the static-scene correspondences used by the temporal alignment methods introduced in Chapter 3. The specific mask generation and integration procedures are described in Chapters 3 and ??.

2.4 Mathematical Background

The reconstruction models evaluated in this thesis do not necessarily predict geometry in a common coordinate system. In particular, the spatial foundation models discussed in Section 2.2 reconstruct each frame or image set independently, often producing point clouds that differ by an arbitrary scale, rotation, and translation. Constructing a temporally coherent 4D reconstruction therefore requires estimating geometric transformations between independently reconstructed frames and combining these estimates into a globally consistent coordinate system.

This section introduces the mathematical concepts that underpin these operations. First, similarity transformations and Umeyama alignment are presented as the basis for point-cloud registration. Next, the geometry of the rotation manifold $SO(3)$ and geodesic rotation averaging are introduced, which form the foundation of the pose-graph optimisation strategy described later in Chapter 3. Finally, the Chamfer distance is presented as a primary metric for evaluating reconstruction quality.

2.4.1 Similarity Transformations and Umeyama Alignment

A three-dimensional similarity transformation aligns two point clouds by accounting for differences in scale, rotation, and translation.

A three-dimensional similarity transformation is defined as

$$\mathbf{y} = s \mathbf{R} \mathbf{x} + \mathbf{t}, \quad (2.1)$$

where $s \in \mathbb{R}^+$ denotes a global scale factor, $\mathbf{R} \in SO(3)$ is a rotation matrix, and $\mathbf{t} \in \mathbb{R}^3$ is a translation vector.

Throughout this thesis, similarity alignment is performed using the least-squares method proposed by Umeyama [20]. Given two sets of corresponding three-dimensional points, the algorithm estimates the optimal scale, rotation, and translation that minimise the mean squared alignment error. Because the solution is available in closed form, Umeyama alignment provides an efficient and numerically stable method for registering reconstructed point clouds.

The estimated similarity transformations are used both for aligning reconstructions to ground truth during evaluation and for computing frame-to-frame alignments in the temporal reconstruction strategies introduced in Chapter 3.

2.4.2 Rotation Geometry on $SO(3)$

While translations and scales can be manipulated using standard Euclidean operations, rotations require special treatment because they lie on a curved manifold rather than in a vector space. The set of valid three-dimensional rotations forms the special orthogonal group

$$SO(3) = \left\{ \mathbf{R} \in \mathbb{R}^{3 \times 3} \mid \mathbf{R}^\top \mathbf{R} = \mathbf{I}, \det(\mathbf{R}) = 1 \right\}. \quad (2.2)$$

Direct arithmetic averaging of rotation matrices generally produces matrices that no longer satisfy the orthogonality constraints of $SO(3)$ [1, 6]. Consequently, optimisation and averaging operations must be performed using geometric constructions defined on the manifold itself.

A commonly used measure of rotational discrepancy is the geodesic distance

$$d(\mathbf{R}_1, \mathbf{R}_2) = \arccos \left(\frac{\text{tr}(\mathbf{R}_1^\top \mathbf{R}_2) - 1}{2} \right), \quad (2.3)$$

which corresponds to the minimum rotation angle required to transform one orientation into the other [10].

2.4.3 Geodesic Rotation Averaging

The pose-graph optimisation strategy developed in this thesis estimates pairwise similarity transformations between many frames in a sequence. Since multiple alignment paths may exist between two frames, individual rotation estimates must be combined into a single globally consistent orientation.

Given a set of rotations $\{\mathbf{R}_i\}_{i=1}^N$, rotation averaging seeks a consensus rotation that best represents the collection. A common formulation minimises the sum of squared geodesic distances,

$$\mathbf{R}^* = \arg \min_{\mathbf{R} \in SO(3)} \sum_{i=1}^N d(\mathbf{R}, \mathbf{R}_i)^2. \quad (2.4)$$

Although effective, this L2 formulation can be sensitive to erroneous pairwise alignments and outliers. To improve robustness, this thesis employs an L1 geodesic formulation,

$$\mathbf{R}^* = \arg \min_{\mathbf{R} \in SO(3)} \sum_{i=1}^N d(\mathbf{R}, \mathbf{R}_i), \quad (2.5)$$

which corresponds to the geometric median on the rotation manifold.

The optimisation is performed using a manifold adaptation of the Weiszfeld algorithm [25]. At each iteration, rotations are mapped to the tangent space of the current estimate via the logarithmic map, robustly averaged using distance-dependent weights, and projected back onto $\text{SO}(3)$ through the exponential map. Compared with conventional L2 averaging, this approach reduces the influence of unreliable pairwise transformations and yields more stable consensus rotations during pose-graph optimisation.

2.4.4 Chamfer Distance

Geometric reconstruction quality is evaluated throughout this thesis using a combination of distance-based and threshold-based metrics. The primary distance-based metric is the symmetric Chamfer distance, which measures the average geometric discrepancy between reconstructed and ground-truth point clouds [2].

Given two point sets P and Q , the Chamfer distance is defined as

$$d_{\text{Chamfer}}(P, Q) = \frac{1}{|P|} \sum_{\mathbf{p} \in P} \min_{\mathbf{q} \in Q} \|\mathbf{p} - \mathbf{q}\|_2 + \frac{1}{|Q|} \sum_{\mathbf{q} \in Q} \min_{\mathbf{p} \in P} \|\mathbf{q} - \mathbf{p}\|_2. \quad (2.6)$$

The first term measures how accurately reconstructed points correspond to the reference geometry, while the second term measures how completely the reference geometry is represented by the reconstruction. Smaller values therefore indicate closer geometric agreement between the two point sets.

Chamfer distance provides a global measure of reconstruction quality and is widely used in three-dimensional vision benchmarks. However, because average distances do not directly indicate whether a reconstruction satisfies the precision requirements of downstream applications, this thesis additionally employs threshold-based Accuracy and Completeness metrics, introduced in Chapter 3, to quantify the proportion of reconstructed geometry lying within a specified error tolerance.

The concepts introduced in this section provide the mathematical foundation for the temporal alignment and evaluation framework developed later in this thesis. In particular, Umeyama alignment enables frame-to-frame registration, while geodesic rotation averaging and Weiszfeld optimisation form the core of the pose-graph optimisation strategy used to construct globally consistent 4D reconstructions from independently reconstructed frames.

3

Methodology: The Evaluation Framework

4

System Implementation

5

Results

6

Refinement Strategy

7

Discussion

8

Conclusions

Bibliography

- [1] Pierre-Antoine Absil, Robert Mahony, and Rodolphe Sepulchre. *Optimization Algorithms on Matrix Manifolds*. Vol. 78. Dec. 2008. ISBN: 978-0-691-13298-3. DOI: [10.1515/9781400830244](https://doi.org/10.1515/9781400830244).
- [2] Harry G. Barrow *et al.* “Image Matching as a Distance Measure”. In: *International Joint Conference on Artificial Intelligence* (1977).
- [3] Yu-Wei Chao *et al.* “DexYCB: A Benchmark for Capturing Hand Grasping of Objects”. In: *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2021.
- [4] Yutong Chen *et al.* “GGPT: Geometry-Grounded Point Transformer”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. 2026.
- [5] Junda Cheng *et al.* “MonSter++: Unified Stereo Matching, Multi-view Stereo, and Real-time Stereo with Monodepth Priors”. In: 2025. URL: <https://api.semanticscholar.org/CorpusID:275544438>.
- [6] Gregory S. Chirikjian. *Stochastic Models, Information Theory, and Lie Groups, Volume 2*. Birkhäuser, 2011.
- [7] Andrea Porfiri Dal Cin *et al.* “AnyMap: Learning a General Camera Model for Structure-from-Motion with Unknown Distortion in Dynamic Scenes”. In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. June 2025, pp. 16674–16684.
- [8] Bardienus Pieter Duisterhof *et al.* “MASt3R-SfM: a Fully-Integrated Solution for Unconstrained Structure-from-Motion”. In: *International Conference on 3D Vision 2025*. 2025. URL: <https://openreview.net/forum?id=5uw1GRBFoT>.
- [9] R. I. Hartley and A. Zisserman. *Multiple View Geometry in Computer Vision*. Second. Cambridge University Press, ISBN: 0521540518, 2004.

- [10] Richard Hartley *et al.* “Rotation Averaging”. In: *International Journal of Computer Vision* 103.3 (2013), pp. 267–305.
- [11] Yu Hu *et al.* *VGGT4D: Mining Motion Cues in Visual Geometry Transformers for 4D Scene Reconstruction*. 2025. arXiv: [2511.19971](https://arxiv.org/abs/2511.19971) [cs.CV]. URL: <https://arxiv.org/abs/2511.19971>.
- [12] Jianfei Jiang *et al.* “MonoMVSNet: Monocular Priors Guided Multi-View Stereo Network”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. Oct. 2025, pp. 27806–27816.
- [13] Vincent Leroy, Yohann Cabon, and Jerome Revaud. *Grounding Image Matching in 3D with MAST3R*. 2024. arXiv: [2406.09756](https://arxiv.org/abs/2406.09756) [cs.CV].
- [14] Wenyu Li *et al.* *Mono3R: Exploiting Monocular Cues for Geometric 3D Reconstruction*. 2025. arXiv: [2504.13419](https://arxiv.org/abs/2504.13419) [cs.CV]. URL: <https://arxiv.org/abs/2504.13419>.
- [15] Zador Pataki *et al.* *MP-SfM: Monocular Surface Priors for Robust Structure-from-Motion*. 2025. arXiv: [2504.20040](https://arxiv.org/abs/2504.20040) [cs.CV]. URL: <https://arxiv.org/abs/2504.20040>.
- [16] Nikhila Ravi *et al.* “SAM 2: Segment Anything in Images and Videos”. In: *arXiv preprint arXiv:2408.00714* (2024). URL: <https://arxiv.org/abs/2408.00714>.
- [17] Johannes Lutz Schönberger and Jan-Michael Frahm. “Structure-from-Motion Revisited”. In: *Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016.
- [18] Johannes Lutz Schönberger *et al.* “Pixelwise View Selection for Unstructured Multi-View Stereo”. In: *European Conference on Computer Vision (ECCV)*. 2016.
- [19] Steven M. Seitz *et al.* “A Comparison and Evaluation of Multi-View Stereo Reconstruction Algorithms”. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (2006).
- [20] S. Umeyama. “Least-squares estimation of transformation parameters between two point patterns”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 13.4 (1991), pp. 376–380. DOI: [10.1109/34.88573](https://doi.org/10.1109/34.88573).
- [21] Jianyuan Wang *et al.* *VGGT: Visual Geometry Grounded Transformer*. 2025. arXiv: [2503.11651](https://arxiv.org/abs/2503.11651) [cs.CV]. URL: <https://arxiv.org/abs/2503.11651>.
- [22] Shuzhe Wang *et al.* “DUST3R: Geometric 3D Vision Made Easy”. In: *CVPR*. 2024.


- [23] Zihan Wang *et al.* “MonoFusion: Sparse-View 4D Reconstruction via Monocular Fusion”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. Oct. 2025, pp. 8252–8263.
- [24] Zihan Wang *et al.* “MonoFusion: Sparse-View 4D Reconstruction via Monocular Fusion”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*. Oct. 2025, pp. 8252–8263.
- [25] Endre Weiszfeld. “Sur le point pour lequel la somme des distances de n points donnés est minimum”. In: *Tohoku Mathematical Journal* 43 (1937), pp. 355–386.
- [26] Yifei Yin *et al.* “Hi4D: 4D Instance Segmentation of Close Human Interaction”. In: *Computer Vision and Pattern Recognition (CVPR)*. 2023.

List of Acronyms



Artifacts Availability

The artifacts associated with this thesis can be found in the ifi-GitLab¹. Within this repository, there are two main folders: one containing the code to fine-tune the model, and the other containing the code to generate the training data, including all the different datasets. The README file provided in the repository contains step-by-step instructions on how to set up the environment used in this thesis, install all necessary dependencies, and run the source code.



Upload your Code to Git-Lab rather than attaching it to the PDF.

¹<https://gitlab.ifi.uzh.ch/ddis/Students/Theses/<NAME>>



Universität
Zürich^{UZH}

Declaration of independence for written work

I hereby declare that I have composed this work independently and without the use of any aids other than those declared (including generative AI such as ChatGPT). I am aware that I take full responsibility for the scientific character of the submitted text myself, even if AI aids were used and declared (after written confirmation by the supervising professor). All passages taken verbatim or in sense from published or unpublished writings are identified as such. The work has not yet been submitted in the same or similar form or in excerpts as part of another examination.

Important: Update this text with the latest **Declaration of independence for written work** from <https://www.ifi.uzh.ch/en/studies/msc-info.html>. Edit me in `resources/declaration_en.tex`

Zürich, June 9, 2026

Signature of
Fabio Di Meo